INDEXER

Zach Scott

The overall Big O of this programming is not terribly large. To go through each token in a given file, the time it would take is O(n), where n is the number of tokens.

The program then needs to add all of these tokens to the hashtable as a new entry, or an existing record. This, also, takes O(n) time for every token encountered during the entire indexing.

Therefore, overall this program will take O(n) time to complete it's job, where N is the number of tokens in ALL the files found during the course of the indexing. This may be further weighted down by Linked-list insertions, which stacks on ADDITIONAL O(n) time for a Record to be inserted at the end of the list, but this is a worst-case scenario.

Hashtable lookup is O(1), which is why I used it :)

This program utilizes UTHASH, created by Troy D. Hanson , to store it's tokens -- all credit where credit is due!